

# Computer Architecture: The Past & The Future

Mingyu Gao

[gaomy@tsinghua.edu.cn](mailto:gaomy@tsinghua.edu.cn)

*IIIS, Tsinghua*

# About Me: Mingyu Gao 高鸣宇

---

- Assistant Professor at Tsinghua University (2019.3. – now)
- Ph.D. from Stanford University (2012.9. – 2018.6.)
- B.S. from Tsinghua University (2008.9. – 2012.7.)
- Research: **Computer Architecture** 计算机系统结构
  - Domain-specific (e.g., AI) systems 领域专用加速
  - Memory systems 存储系统架构
  - Hardware security 硬件安全计算
  - .....
- More on my webpage: <http://people.iis.tsinghua.edu.cn/~gaomy/>

# Terminology & Vocabulary

---

- Von Neumann Architecture 冯诺依曼架构
- Arithmetic/logic unit (ALU) 算术逻辑单元
- Register file 寄存器堆
- Memory 内存
- Assembly code 汇编代码
- Machine code 机器码
- Instruction 指令
- Instruction set architecture (ISA) 指令集架构
- Moore's Law 摩尔定律
- Integrated Circuits 集成电路
- Transistor 晶体管
- Parallelism 并行
- Pipelining 流水线执行
- Superscalar 超标量
- Out-of-order 乱序执行 Vector 矢量
- Multi-thread 多线程
- Bandwidth 带宽

# COMPUTER SYSTEM BASICS

How a computer works, in a *good* way

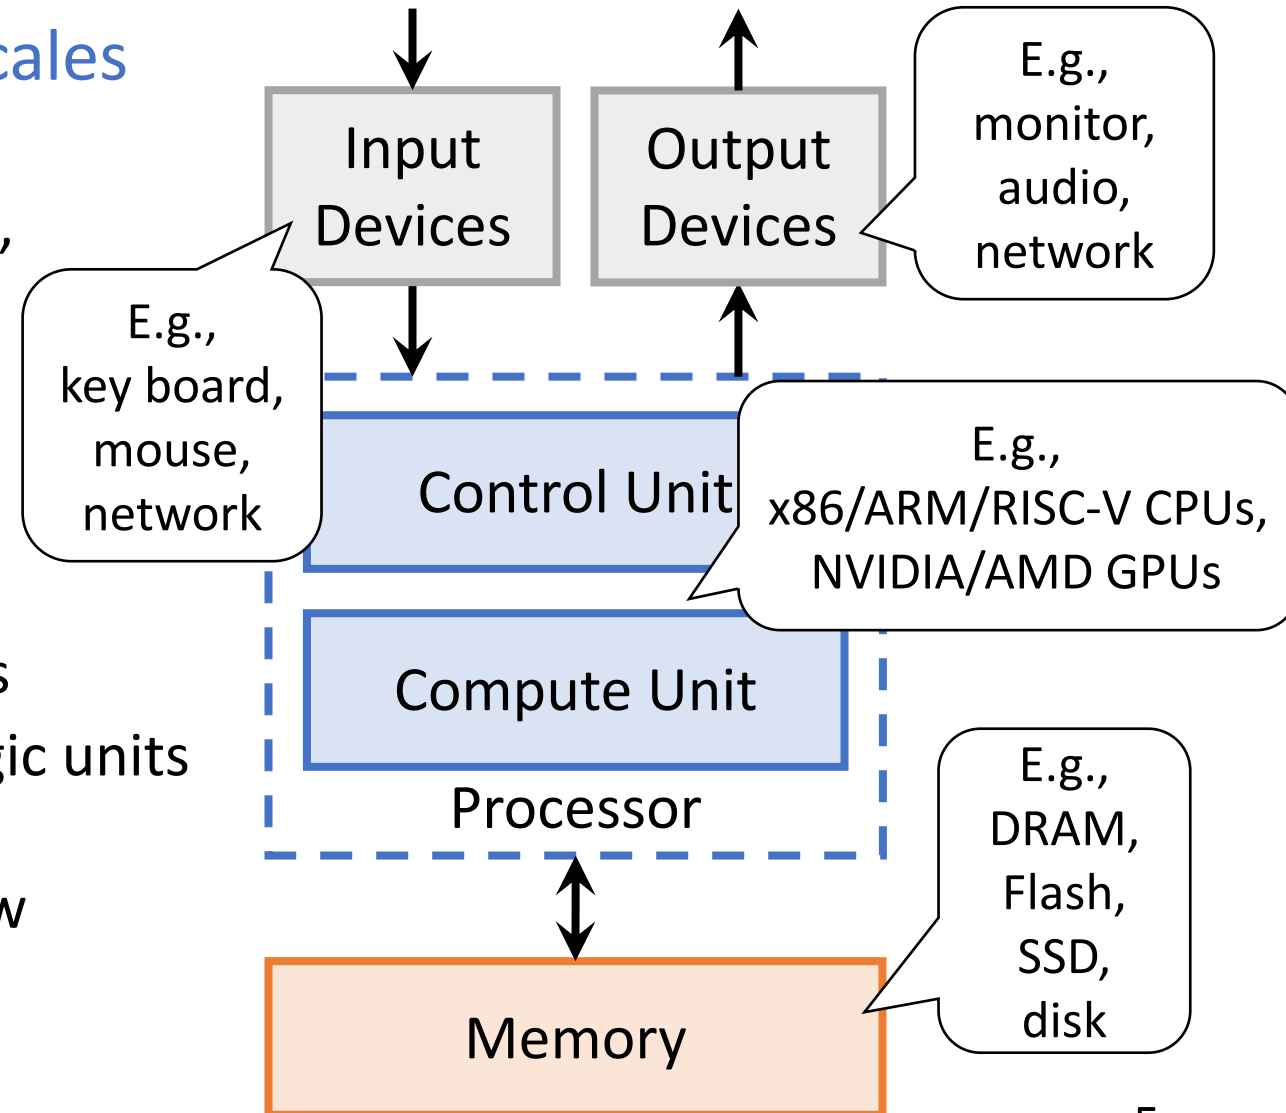
# Typical Computer Systems Today

- Computer systems have **diverse scales** and **abstraction levels**

- Mobile phones, laptops, desktop PCs, distributed clusters, datacenters, ...
- We study all of them

- Von Neumann Architecture**

- Memory**: for both data and programs
- Compute** unit: include arithmetic/logic units and registers
- Control** unit: coordinate program flow
- Input/output** devices

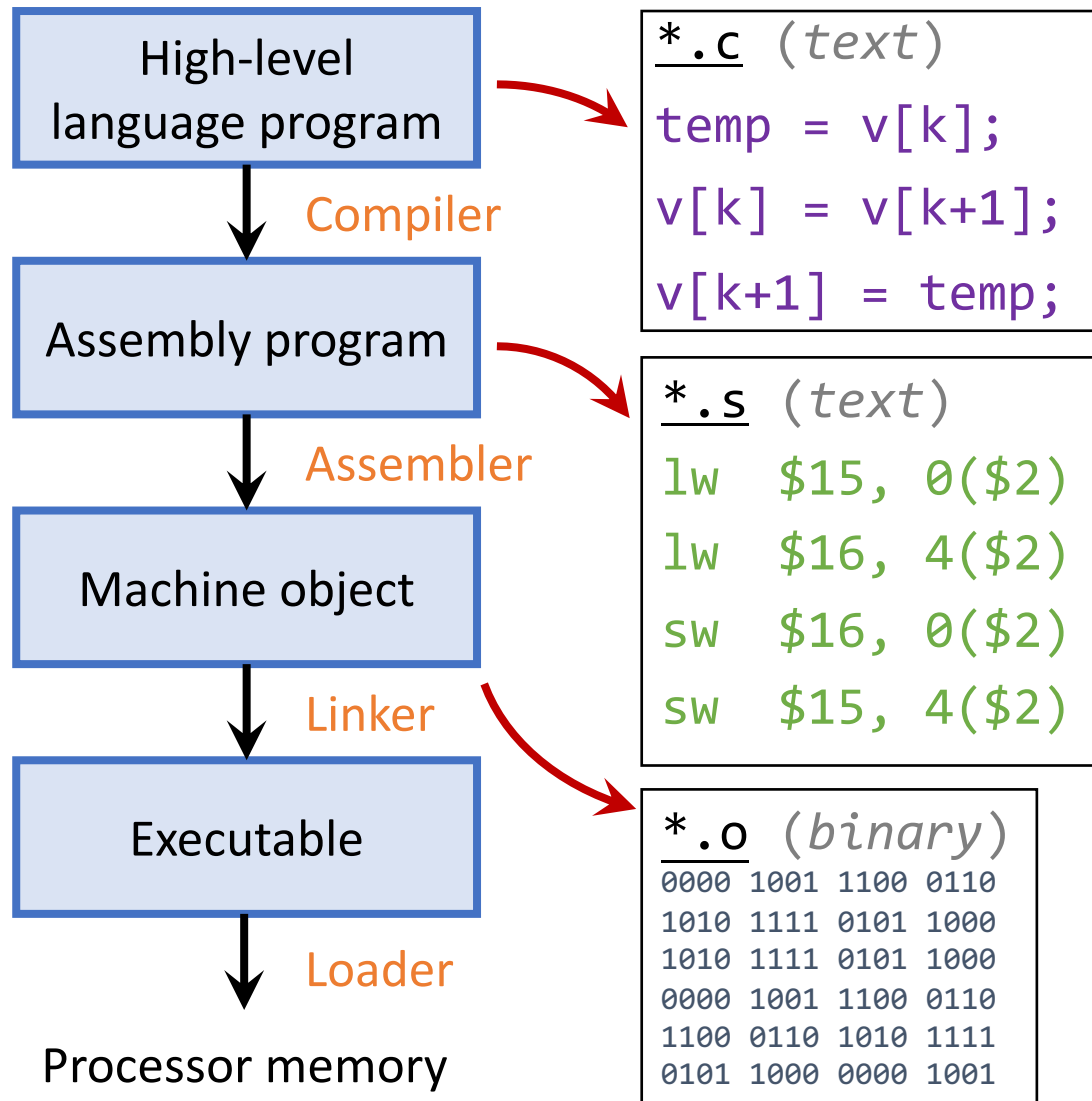


# How Computer Systems Work

---

- ❑ Computers work with binary signals, i.e., bits (0 and 1)
  - Everything is expressed as sequences of bits
- ❑ **Program**: sequence of **instructions**, encoded as strings of bits
  - ISA (instruction set architecture)
  - Example ISAs: x86, x86-64, ARM, RISC-V, MIPS, ...
- ❑ **Data storage** (i.e., **memory**): cells preserve bits over time
  - Flip-flops, registers, SRAM, DRAM, ...
- ❑ **Data processing** (i.e., **compute & control**): logic gates operate on bits
  - AND, OR, NOT, MUX, add, mult, ...

# Code Translation

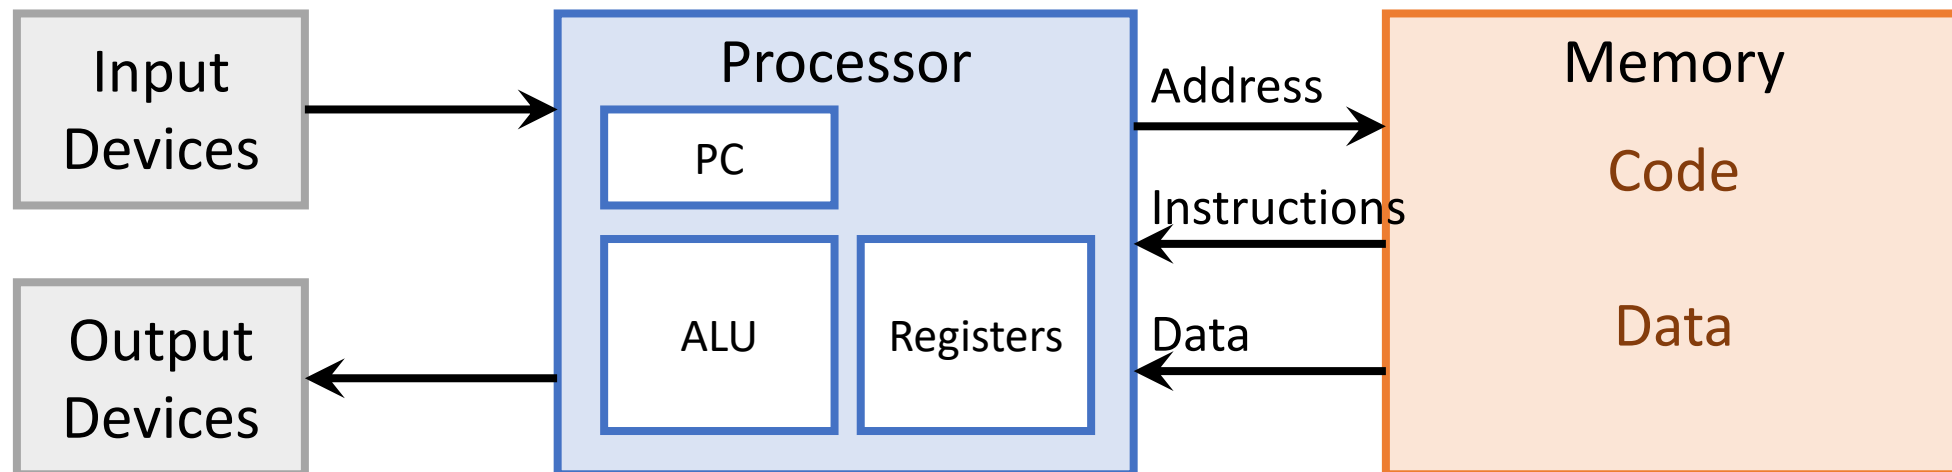


- ❑ **Machine code**: the byte sequence that encodes program instructions
  - Actual 0's and 1's stored in the memory
- ❑ **Assembly code**: text representation of machine code
  - Human-readable instruction sequence
- ❑ **Assembler**
  - Binary encode each instruction
- ❑ **Linker**
  - Resolve inter-file references
  - Combine with static libraries
    - Some libraries are dynamically linked

# Assembly's (Simplified) View of The System

---

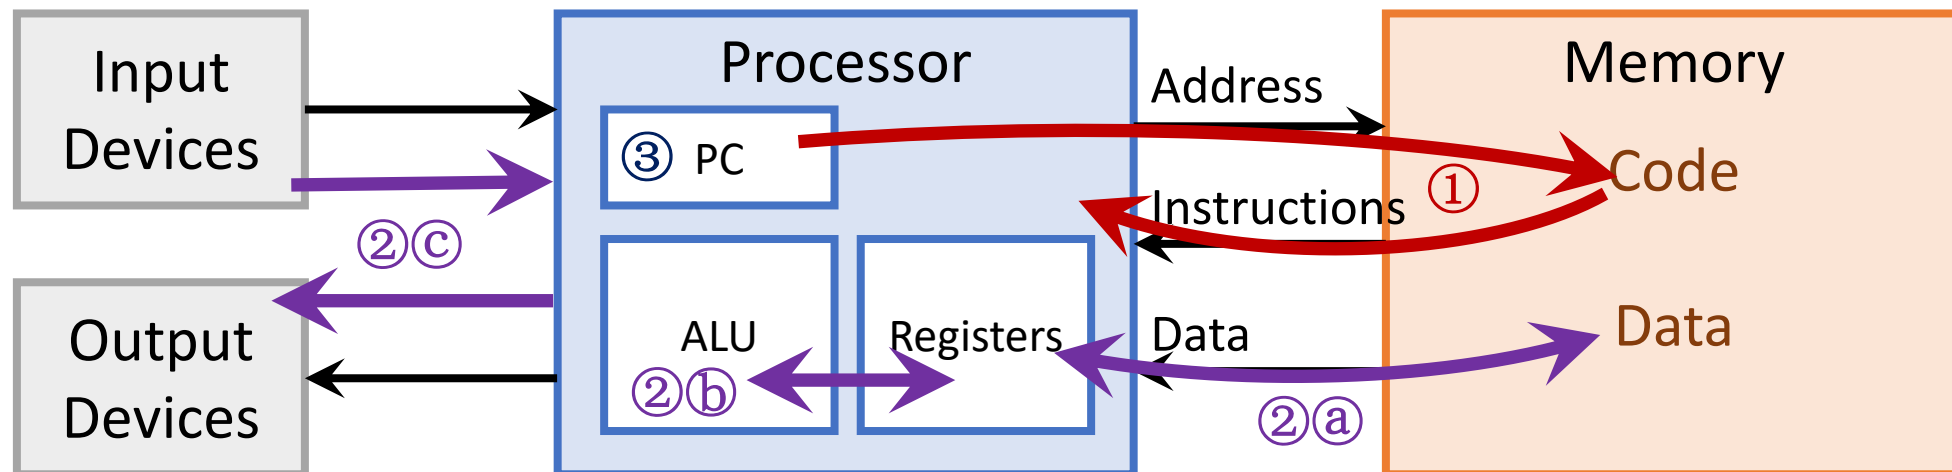
- ❑ State: PC, registers, memory
  - Program counter (PC): points to the next instruction
    - Where we are in the middle of the instruction sequence
  - Registers (a.k.a., register file): local, heavily used data in the processor
  - Memory: code and data, both stored as blocks of bytes
- ❑ Arithmetic/logic unit (ALU): digital logic to do computation





# How Instructions Execute

- 1) PC used as address to fetch an instruction from memory
- 2) The instruction indicates how to compute
  - a) Transfer data between memory and processor registers
  - b) Compute on register values and store results to registers
  - c) Special ways to interact with outside world through I/O
- 3) Update PC to the next instruction, go back to 1)



# What We Always Want

---

- All computer systems should work fast & efficiently
- High performance
  - Theorists invent algorithms by considering asymptotic behaviors (e.g.,  $O(n \log n)$ )
  - Architects set the constant factors
- Low cost, in terms of money, area, power, energy, design complexity, ...
  - It is not hard to have high performance, but it is hard to do so with low cost

*How did we achieve this goal before?*

*What could we do next to continue?*

# THE OLD GOOD DAYS

... when we have everything *scale* well



# Moore's Law

- Gordon E. Moore
  - Co-founder of Fairchild, later became CEO of Intel
- The original statement (1965)

*The complexity for minimum component costs has increased at a rate of roughly a factor of two per year*
- The more well-known version

*The number of transistors per square inch on integrated circuits is doubling every year*

  - In 1975, revised to *doubling every two years*





Fig. 2 Number of components per integrated function for minimum cost per component extrapolated vs time.

# How Is Moore's Law Doing?

## Moore's Law – The number of transistors on integrated circuit chips (1971-2018)

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important as other aspects of technological progress – such as processing speed or the price of electronic products – are linked to Moore's law.



- ❑ 1974, Intel 8080
  - 6,000 transistors
  - 20 mm<sup>2</sup>
- ❑ 2018, HiSilicon Kirin 980
  - 6,900,000,000 transistors
  - 74.13 mm<sup>2</sup>
- ❑ 310,000× in 44 years
  - 1.33× per year
  - About 1.7× per two years
- ❑ Moore's Law is still alive!


Data source: Wikipedia ([https://en.wikipedia.org/wiki/Transistor\\_count](https://en.wikipedia.org/wiki/Transistor_count))  
The data visualization is available at [OurWorldinData.org](https://ourworldindata.org). There you find more visualizations and research on this topic.

Licensed under [CC-BY-SA](#) by the author Max Roser.

# What is Happening Under the Hood

## □ Transistor

- MOS: Metal-Oxide-Semiconductor
- Digital switch: G controls the current from S to D
- Use these 0/1 switches to build binary circuits



## □ Feature size: length of the transistor channel

- Feature size shrinks in each technology generation (roughly  $1.4\times$  reduction)
  - $0.18\text{ }\mu\text{m} \rightarrow 0.13\text{ }\mu\text{m} \rightarrow 90\text{ nm} \rightarrow 65\text{ nm} \rightarrow 45\text{ nm} \rightarrow 32\text{ nm} \rightarrow 22\text{ nm} \rightarrow 14\text{ nm} \rightarrow 10\text{ nm} \rightarrow 7\text{ nm} \rightarrow \dots$
- Other dimensions shrink along with feature size with roughly same ratio
  - For planar area,  $2\times$  reduction per generation, enabling Moore's Law

# Dennard Scaling

---

- ❑ Moore's Law is only about chip density scaling
  - If the transistor feature size scales by  $1/S$  ( $S \approx 1.4$ ), the chip density scales by  $S^2$
- ❑ Performance and cost: **Dennard Scaling** (Robert H. Dennard, 1974)
  - If the feature size scales by  $1/S$ , the supply voltage and current can scale by  $1/S$
- ❑ The free lunch in semiconductor industry!
  - **2.8x capability** (2x transistors & 1.4x speed) with **same area** and **same power**
  - E.g., Intel Tick-Tock model

Feature Size	Area $A$	Capacitance $C$	Voltage $V$	Current $I$	Freq $f = I/CV$	Energy $E = CV^2$	Power $P = CV^2 f$	Pwr Density $P/A$
$1/S$	$1/S^2$	$1/S$	$1/S$	$1/S$	$S$	$1/S^3$	$1/S^2$	1

# From More Transistors to Higher Performance

---

- ❑ Processor architectures have been evolving to utilize the increasing amount of transistors to offer higher performance
- ❑ By leveraging different levels of **parallelism**
  - ❑ Bit-level parallelism
  - ❑ Instruction-level parallelism
  - ❑ Data-level parallelism
  - ❑ Thread-level parallelism



# Bit-Level Parallelism

---


- ❑ From 1970s to mid 1980s, increase datapath bit width
  - 8-bit: 0 to 255, or -128 to 127
  - 16-bit: 0 to 65536, or -32768 to 32767
  - 32-bit: 0 to 4G, or -2G to 2G, large enough in most cases, except for memory
  - 64-bit: little benefit to further increase

Processor	Year	# Transistors	Bit width
Intel 4004	1971	2,250	4-bit
Intel 8008	1972	3,500	8-bit
Intel 8086	1978	29,000	16-bit
Intel 80386	1985	275,000	32-bit


- ❑ Reduce number of cycles required to perform arithmetic operations
  - E.g., multiply two 32-bit numbers on a 16-bit processor
    - Four multiplies ( $A_H \times B_H, A_H \times B_L, A_L \times B_H, A_L \times B_L$ ), followed by several adds

# Instruction-Level Parallelism

- ❑ From mid 1980s to 2000, increase **instruction throughput**
  - **Pipelining**: partially overlap execution of instructions
  - **Superscalar & out-of-order**: executing multiple instructions together
- ❑ **Limitations**
  - **Dependencies (data or control)** between instructions
  - **Hardware complexity and cost**



Pipelining




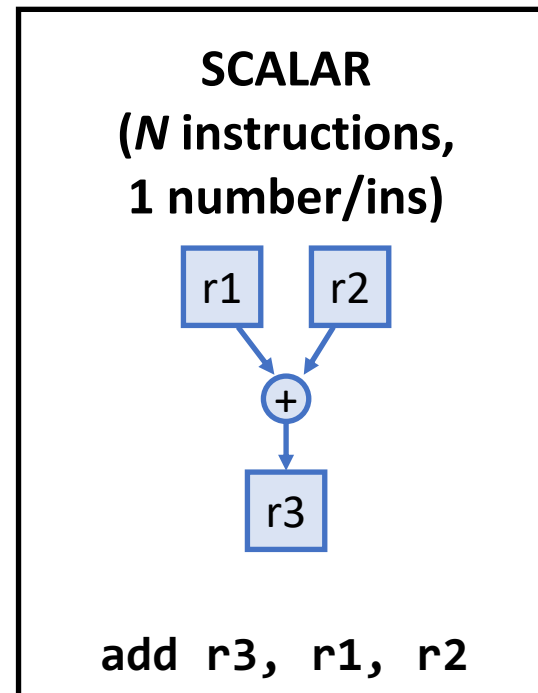
Superscalar & OoO

# Data-Level Parallelism

- ❑ Scalar processors execute each instruction on single numbers (scalars)
- ❑ Vector processors may execute one instruction on vectors of numbers
- ❑ **SIMD** (single-instruction, multiple-data)
  - Increase processing throughput
  - Amortize instruction cost
  - E.g., Intel SSE, AVX, ...


```
for i: 0 to N  
    c[i] = a[i] + b[i]
```

Q: what domains of algorithms  
would SIMD be most beneficial?



# Thread-Level Parallelism

- ❑ After 2000, **chip multiprocessors (CMPs)**, i.e., **multi-core**
- ❑ Single-thread performance is more and more difficult to optimize
- ❑ Explicitly extract task parallelism and specify as multi-thread programs



# THE WALLS

... which block us from further improving performance *now*



# Post-Dennard Scaling Era

---

- ❑ Moore's Law without Dennard Scaling
  - No more voltage and current scaling/reduction!! Why?
- ❑ When transistor size is so small (a few nanometers) ...
  - Transistor threshold voltage and gate oxide thickness are set by leakage
  - Further reducing voltage cannot effectively turn off
  - Leakage power may exceed switch power
- ❑ Implication: 1.4x chip capability per generation if with constant power

Feature Size	Area $A$	Capacitance $C$	Voltage $V$	Current $I$	Freq $f = I/CV$	Energy $E = CV^2$	Power $P = CV^2 f$	Pwr Density $P/A$
$1/S$	$1/S^2$	$1/S$	$\sim 1$	$\sim 1$	$< S$	$1/S$	$< 1$	$< S^2$

# The Power Wall

---

- ❑ Power Wall: all computers are power limited
  - From cellphones to datacenters
  - Power density would increase at  $S^2$ , need to cap it at 1
- ❑ Slow down the frequency scaling
  - Basically frequency has almost stopped scaling (see previous slide):  $S \rightarrow 1$
- ❑ Reduce chip utilization (*Dark Silicon*: dark area of chip not turned on)
  - $S^2$  transistors available, but can only utilize  $S$  transistors simultaneously
  - Save another  $S$  factor

# Recall: Von Neumann Architecture

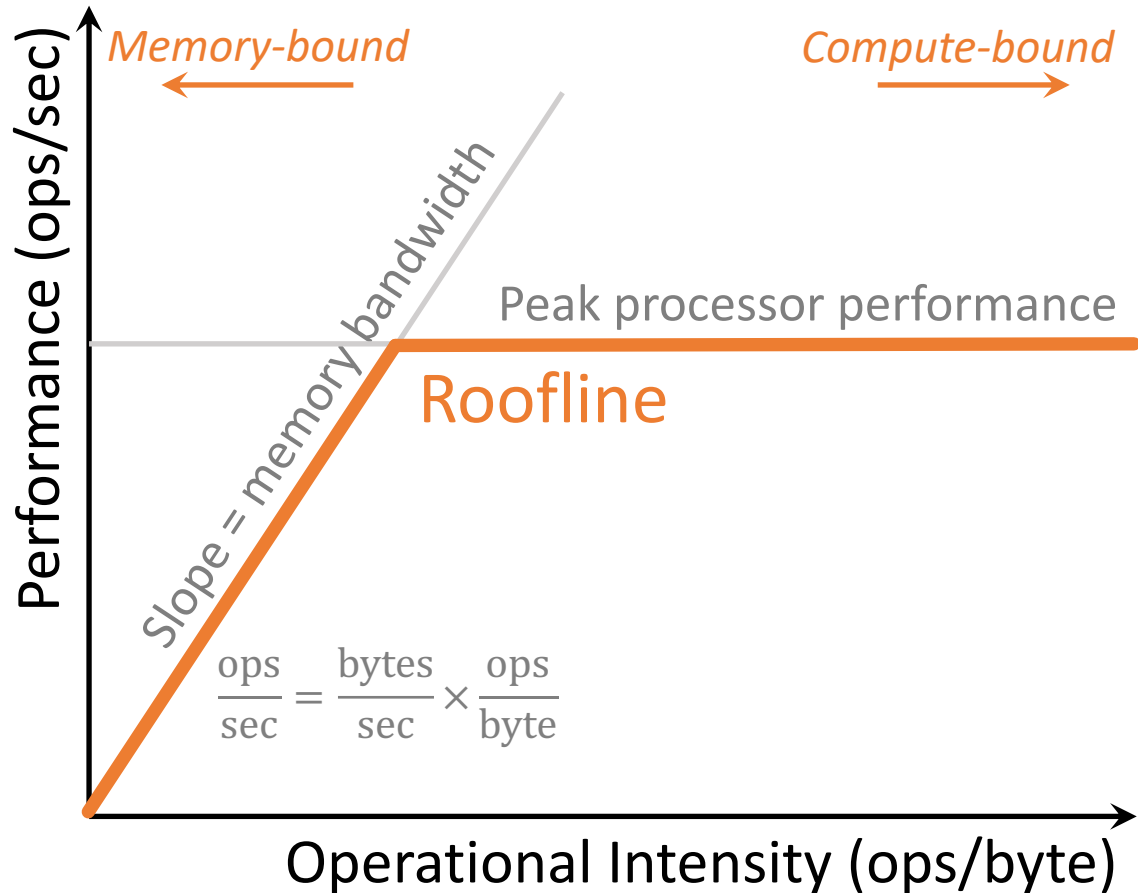
---

- ❑ Separated **processor** (compute chip) and **memory** (data chip)
  - Fetch data from memory into processor
  - Compute in processor (processor has limited local data store)
  - Write back data from processor to memory
  - Continue with next data ...
- ❑ Overall performance is determined by both processor and memory
  - How fast processor computes
  - How fast memory delivers data







# The Roofline Model



- Performance: operations per second
- Operational Intensity
  - How many ops to perform for each byte
  - A characteristic of program
  - Denote the ratio of compute to data
- Overall performance = min(  
processor performance,  
memory performance)
  - High OI: compute-bound
  - Low OI: memory-bound

# The Memory Wall

- ❑ Memory Wall: memory performance and energy dominate systems
  - Memory (DRAM) technology scales more slowly than processor (MOS)
  - Today, memory accesses are **two to three orders of magnitude** more expensive than processor operations



# MAKE THINGS GREAT AGAIN

Let's break the walls!




# Improving Power: Energy Efficiency

$$\text{Power} = \frac{\text{Energy}}{\text{Time}} = \frac{\text{Energy}}{\text{Op}} \times \frac{\text{Op}}{\text{Time}}$$

Energy Efficiency                      Performance

- To improve performance we must improve **energy efficiency**
  - Otherwise, we cannot use the additional transistors
- Where does the energy go?

**Figure 4. Processor energy breakdown for base implementation.** IF is instruction fetch/decode. D-\$ is data cache. P Reg includes the pipeline registers, buses, and clocking. Ctrl is miscellaneous control. RF is register file. FU is the functional units.



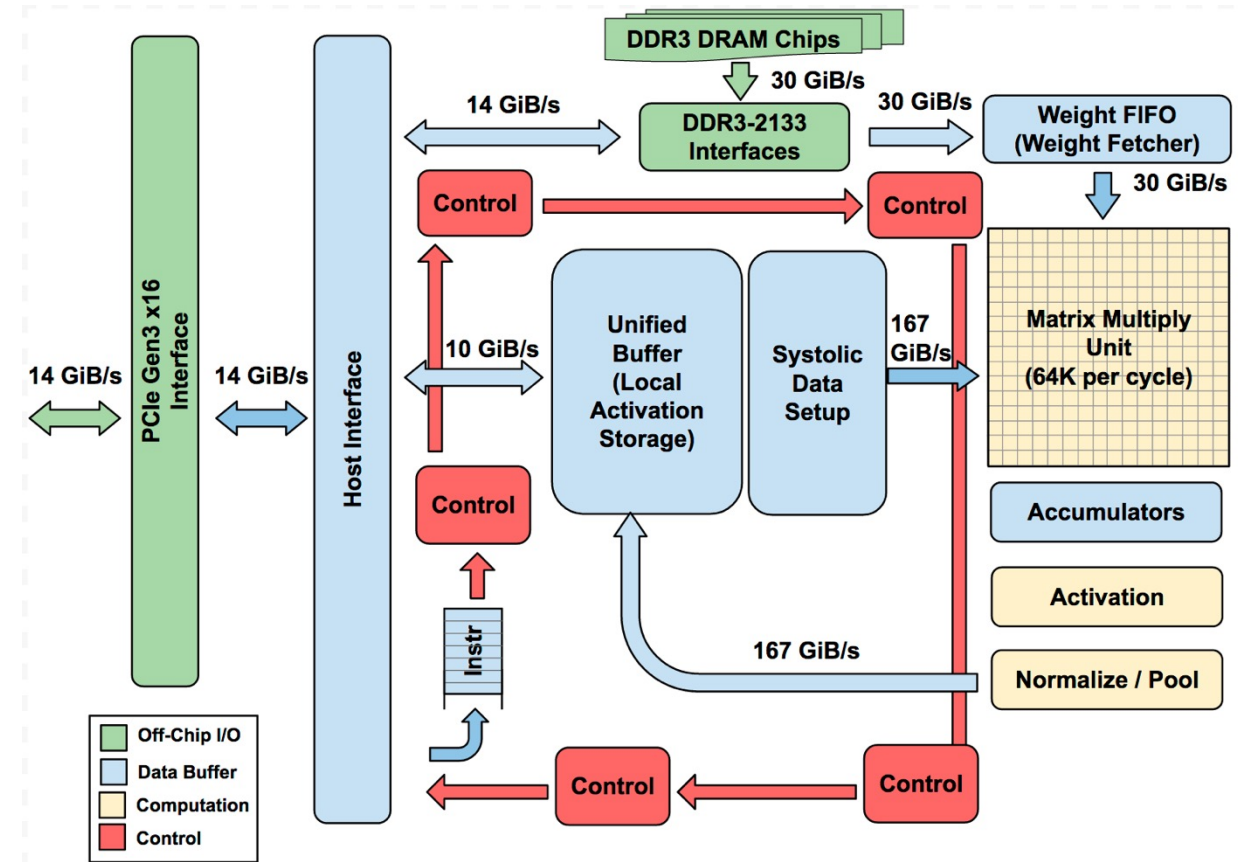
# Domain-Specific Architectures (DSAs)

---

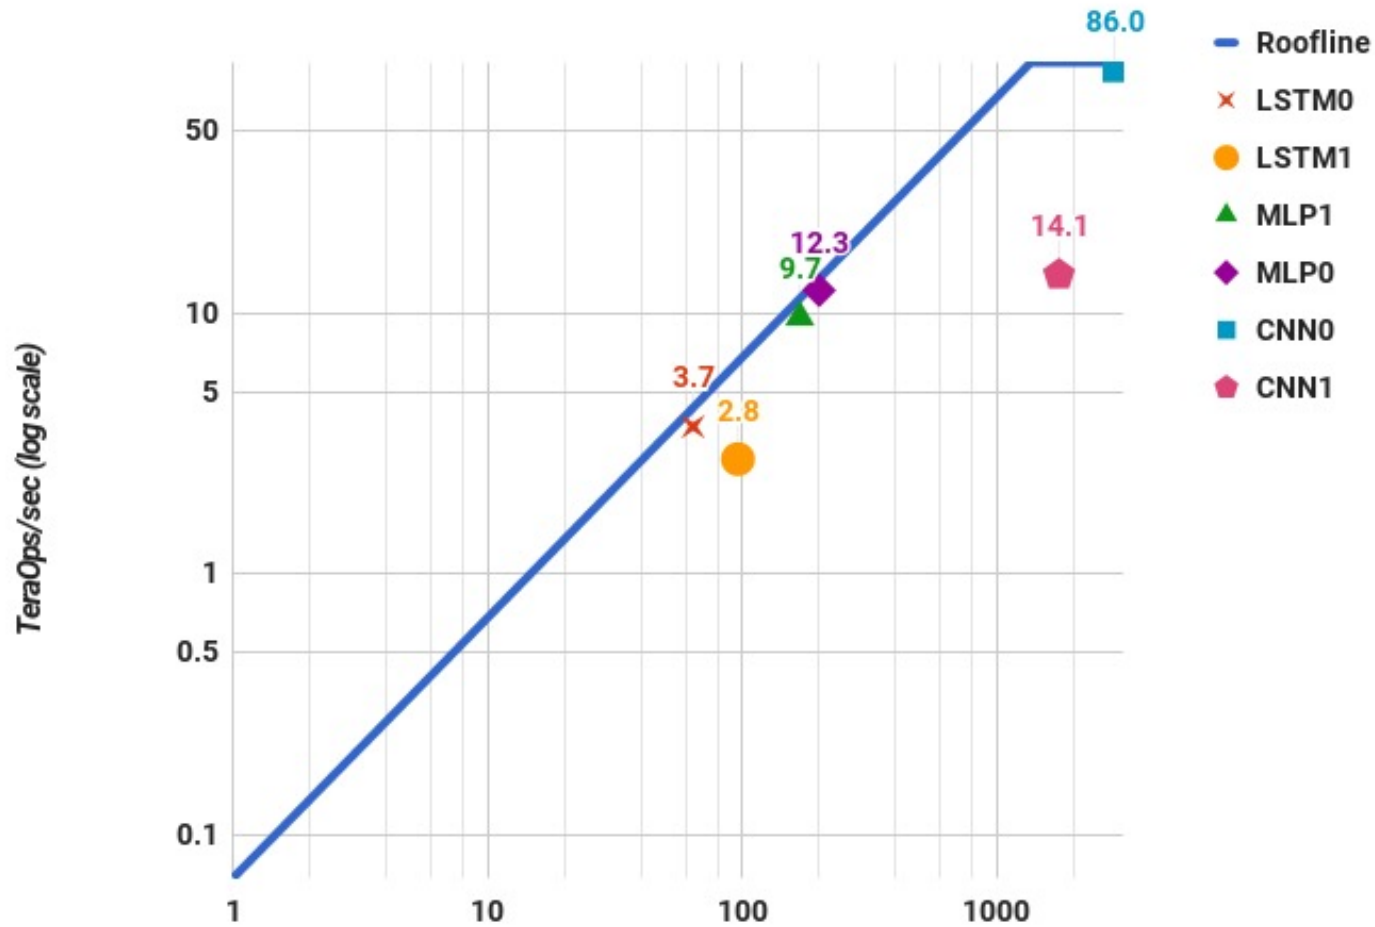
- ❑ Achieve better energy efficiency by **tailoring the architecture to characteristics of the domain**
  - Not one application, but a domain of applications – different from strict ASIC
  - Requires more domain-specific knowledge than general-purpose processors
- ❑ DSA can ...
  - **Reduce overheads of general-purpose**
    - Instruction fetch, instruction scheduling (pipelining/OoO), ...
  - **Use custom optimizations**
    - Lower precision in certain domains, e.g., neural nets
  - **Have more effective parallelism**
    - SIMD, spatial arrays of simple processing elements, ...
  - **Access memory more efficiently**
    - User-controlled on-chip buffer management

# Example: Google TPU v1 for Neural Nets

- Matrix Multiply Unit
  - 256 x 256 8-bit multiply-accumulate
  - Peak: 92 Tops/sec (@ 700 MHz)
- On-chip memories
  - 4 MB accumulator memory
  - 24 MB unified buffer
- 25x MACs vs. GPU
- 3.5x on-chip memory vs. GPU
- 29x perf/Watt vs. GPU



# TPU v1 Roofline



□ What do you see from this plot?

Operational Intensity: MAC Ops/weight byte (log scale)

# Open Research Questions for DSA


---

- ❑ Tradeoff between general-purpose and domain-specific
  - Generality, cost vs. performance, efficiency
  - Smart selection of domains
- ❑ Hardware development cost → agile hardware development
  - Debugging, testing, verification, ...
- ❑ Programming on DSAs
  - Domain-specific languages: TensorFlow, OpenGL, P4, Halide, ...
  - Compiler challenges: from DSL code to DSA binary
  - Co-design of new DSLs and DSAs




# Improving Memory

- ❑ Memory **capacity** is not too hard to scale up
- ❑ Data bandwidth to processor is limited, restricting **performance**
  - $\text{Bandwidth (bytes/sec)} = \text{datawidth (bytes)} \times \text{frequency (Hz)}$
  - Channel datawidth is limited by available chip pins, which is limited by chip area
  - Channel data frequency is limited by signal integrity
- ❑ Inter-chip data transfers (memory to processor) have high **energy** cost

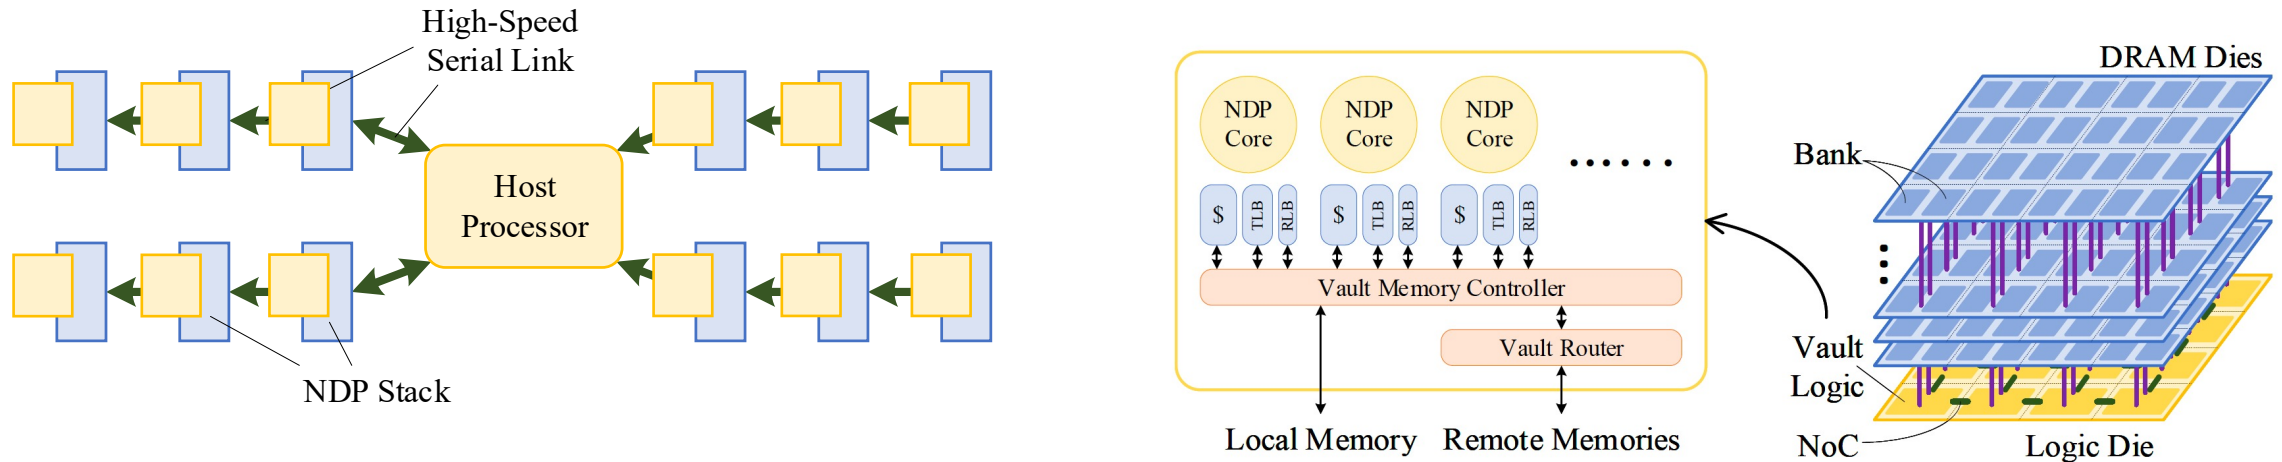


# Near-Data Processing/Processing-In-Memory

- ❑ Key idea of NDP/PIM: **execute computation closer to data**
- ❑ Attach small processors inside or near memory modules
  - **High local data bandwidth, low energy cost, low access latency, high parallelism**
- ❑ Near-data compute capability is constrained by **area & power budget**
  - Compute-bound tasks: execute on host processor
  - Memory-bound tasks: execute on near-data processors

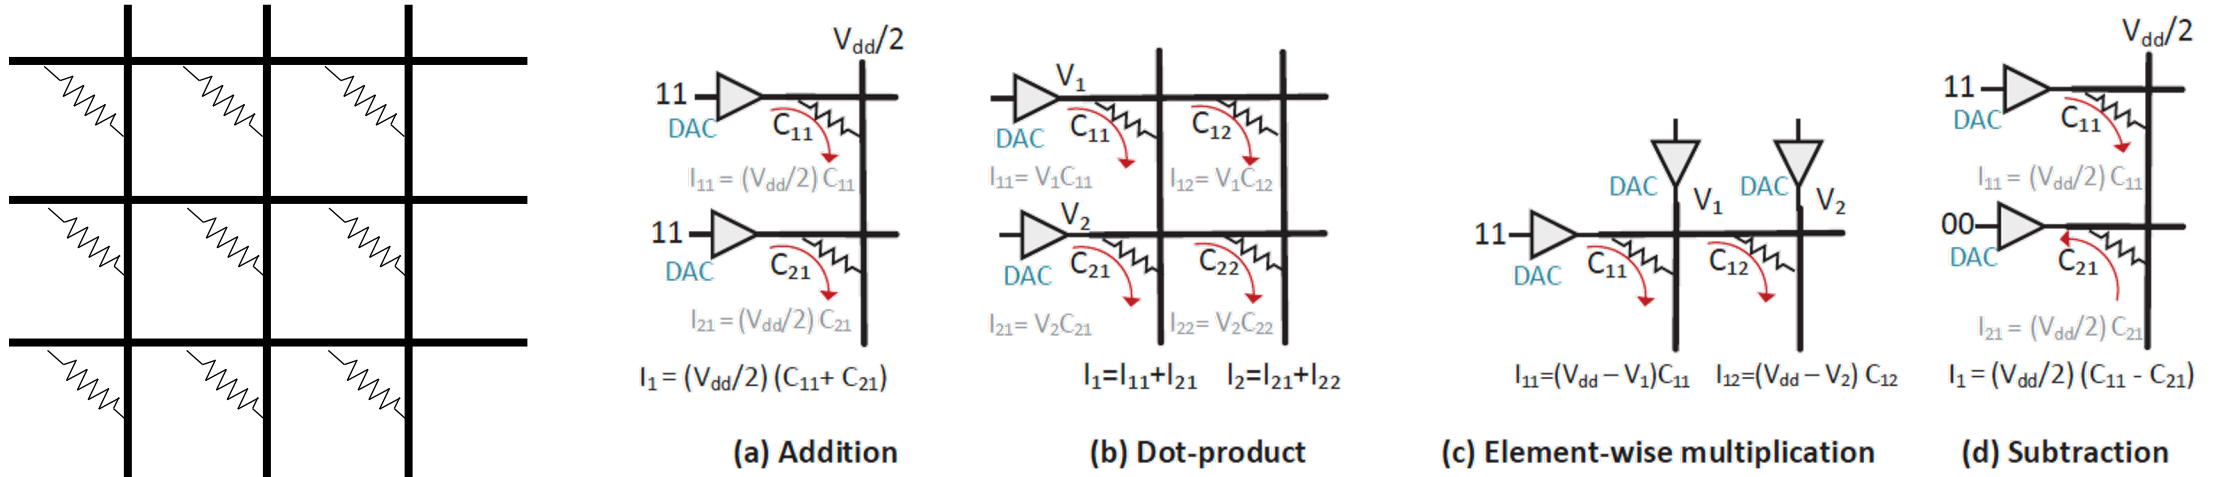


# Example: Near-Data Processing Architecture



- ❑ **3D stacking integration** for NDP: a logic die below multiple memory dies
  - Much more local vertical connections within a chip module
  - Memory bandwidth scale proportionally to memory capacity
  - Much more small cores utilize abundant memory bandwidth
- ❑ **Implication: centralized → distributed**
  - Making software programming more difficult

# Example: In-Situ Analog Processing-in-Memory



- ❑ **Resistive RAM (memristor):** cell resistance/conductance represent value
- ❑ Enable **analog dot-product** ( $I = \sum V \cdot G$ ) and more
  - Computation within memory: one of the two operands stays inside memory
  - Require digital-to-analog and analog-to-digital conversions (DAC, ADC)
  - Well-suited for vector/matrix computation
- ❑ Challenge: how to enable more general, fine-grained operations?

# Open Research Questions for NDP/PIM

---

## □ Impact on programming models

- Heterogenous between host and NDP processors
- Distributed among NDP processors
- Work partition, data layout, coherence, consistency, synchronization, ...

## □ Issues with analog computation

- Overheads of digital-analog conversion
- Physical device reliability, noise, interference, data precision, ...

## □ More efficient computation

- Fully utilize abundant memory bandwidth
- Meet the tight constraints of area and power budget in memory modules

# Summary

---

- ❑ Computer systems used to be great! – exponential perf increasing
  - Moore's Law and Dennard Scaling
  - Parallelism in processor architectures
- ❑ They are not so “great” now ☹ – the challenges
  - *The Power Wall*
  - *The Memory Wall*
- ❑ How to make computer systems great again? – the research directions
  - Domain-specific architectures
  - Near-data processing/processing-in-memory
  - And many more!

# The Scope of My Research Lab

---

## *Innovative Data-centric Efficient Architecture Lab (IDEAL)*



### □ Data storage

- Near-Data Processing/Processing-in-Memory architectures
- Hybrid memory systems: DRAM + NVM



### □ Data processing

- AI accelerators: hardware architectures + software scheduling
- Reconfigurable architectures



### □ Data security

- Isolated execution on accelerator hardware
- Hardware enclaves + cryptographic algorithms

# Extended Reading/Watching

---

- ❑ Turing Lecture by 2017 ACM A. M. Turing Award recipients – John L. Hennessy and David A. Patterson

*A New Golden Age for Computer Architecture:  
Domain-Specific Hardware/Software Co-Design,  
Enhanced Security, Open Instruction Sets,  
and Agile Chip Development*

- ❑ <https://dl.acm.org/citation.cfm?id=3282307>
  - A video recording the Turing lecture (on YouTube)
  - An article on Communications of the ACM



# THANKS!

Webpage: <http://people.iis.tsinghua.edu.cn/~gaomy/>

Contact: [gaomy@tsinghua.edu.cn](mailto:gaomy@tsinghua.edu.cn)